

HAOKUN CHEN

Munich, Germany ◊ (+49) 15237770369 ◊ chenhaokun24549@gmail.com ◊ haokunchen245.github.io

EDUCATION

Ph.D. Computer Science , LMU Munich, Germany	<i>11.2021 - 10.2025</i>
Interests: Model Post-Training : Preference Optimization & Alignment of LLM. Distributed Training : Privacy-preserving Federated Learning. Trustworthy AI : LLM/VLM Red-Teaming, Machine Unlearning, Adversarial Robustness.	
M.Sc. Computer Science , Technical University of Munich, Germany	<i>10.2018 - 10.2021</i>
B.Sc. Mechatronic Engineering , Tongji University, China	<i>10.2014 - 10.2018</i>

WORK EXPERIENCE

Applied Scientist Intern , Amazon Web Services (Berlin, Germany)	<i>06.2025 - 12.2025</i>
○ Researched diversity-enhanced post-training for LLM to improve product review summarization. ○ Developed a voice-based review creation assistant using LLM agents for content, title, and star rating suggestions, integrated with Amazon Bedrock Guardrails to ensure compliant generation.	
Research Scientist Intern , Intel Labs (US, Remote)	<i>08.2024 - 11.2024</i>
○ Designed and implemented an activation-steering algorithm and an evaluation framework to systematically audit existing machine unlearning methods for LLMs. ○ Contributed to LLM adversarial robustness toolkit (LLMart): https://github.com/IntelLabs/LLMart to facilitate the evaluation of adversarial attack algorithms for LLMs.	
Machine Learning Researcher , Siemens AI Lab (Munich, Germany)	<i>01.2021 - 06.2025</i>
○ Implemented a scalable training pipeline (across 16 nodes) for CAD drawing embedding models within RAG-based industrial chatbot on Microsoft Azure ML. ○ Built a hyperparameter tuning framework for LLM post-training across multi-node distributed systems, reducing training time by 15%; work accepted at <i>AAAI 2025</i> (Oral, First Author). ○ Proposed the first PEFT-based post-training algorithm for Vision-Language Models (VLMs) within privacy-preserving federated learning across heterogeneous distributed datasets; manuscript accepted at <i>AAAI 2024</i> (First Author). ○ Designed algorithms to address data heterogeneity, scarcity, and system scalability challenges in federated learning; manuscripts accepted at <i>ICCV 2023</i> and <i>CVPR 2025</i> (Both First Author). ○ Conducted research on adversarial robustness and red-teaming strategies for Multimodal LLMs. ○ Contributed to customer-focused data science and ML projects across various business units. ○ Provided technical leadership and mentorship to master's students and junior PhD researchers.	
Research Intern , BMW Autonomous Driving Campus (Munich, Germany)	<i>03.2020 - 11.2020</i>
○ Developed validation frameworks for computer vision functions in autonomous driving (C++). ○ Conducted research on data augmentation techniques for traffic sign detection.	
Software Engineer Intern , BMW Research Center (Beijing, China)	<i>08.2019 - 12.2019</i>
○ Contributed to the development of validation frameworks for in-vehicle entertainment systems.	

PUBLICATIONS

Haokun Chen, et al., AUVIC: Adversarial Unlearning of Visual Concepts for Multi-Modal Large Language Models, **AAAI 2026**.

Haokun Chen, et al., FedBiP: Heterogeneous One-Shot Federated Learning with Personalized Latent Diffusion Models, **CVPR 2025**.

Haokun Chen, et al., Soft Token Attacks Cannot Reliably Audit Unlearning in Large Language Models, **EMNLP (Findings) 2025**.

Haokun Chen, et al., FedPop: Federated Population-based Hyperparameter Tuning, **AAAI 2025, Oral**.

Haokun Chen, et al., FedDAT: An Approach for Foundation Model Finetuning in Multi-Modal Heterogeneous Federated Learning, **AAAI 2024**.

Haokun Chen, et al., FRAug: Tackling Federated Learning with Non-IID Features via Representation Augmentation, **ICCV 2023**.

Yao Zhang, **Haokun Chen**, et al., CI-crossvqa: A continual learning benchmark for cross-domain visual question answering, **WACV 2025**.

Jinhe Bi, Yujun Wang, **Haokun Chen**, et al., LLaVA Steering: Visual Instruction Tuning with 500x Fewer Parameters through Modality Linear Representation-Steering, **ACL (Main) 2025**.

Yao Zhang, Chenyang Lin, **Haokun Chen**, et al., SwarmAgentic: Towards Fully Automated Agentic System Generation via Swarm Intelligence, **EMNLP (Main) 2025**.

Haokun Chen, et al., Does Gradient Ascent Truly Remove Model Knowledge? A Framework for Auditing Unlearning in LLMs, *Under review*, 2025.

Shuo Chen, Zhen Han, **Haokun Chen**, et al., Bag of Tricks for Subverting Reasoning-based Safety Guardrails, *Under review*, 2025.

Yao Zhang, Hewei Gao, **Haokun Chen**, et al., FedNano: Toward Lightweight Federated Tuning for Pretrained Multimodal Large Language Models, *Under review*, 2025.

Ahmed Frikha*, **Haokun Chen***, et al., Towards Data-Free Domain Generalization, **ACML 2022**.

SKILLS

Programming Languages: Python, C++, Unix, Git.

Frameworks: PyTorch, Accelerate, DeepSpeed, Ray-Tune, TensorRT-LLM, Docker, vLLM.

Languages: English, Mandarin, German (C1).

REVIEW

Conferences: *Computer Vision:* CVPR 2025/2026, ECCV 2024, ICCV 2025.

Machine Learning: NeurIPS 2024/2025, ICLR 2025/2026, ICML 2025.

Artificial Intelligence: AAAI 2024/2025/2026, AISTATS 2025/2026, IJCNN 2025.

Journals: TPAMI, TKDE, TKDD, TCNN